

RESEARCH INTERNSHIP REPORT
SOFTWARE SCIENCE



RADBOD UNIVERSITY

**ChatGPT's role in developer
conversations across software
platform**

Author:
Elina Antonova
s1057069

Supervisor:
dr. Mairieli Wessel
mairieli.wessel@ru.nl

July 7, 2025

Abstract

This study investigates the role of Chat Generative Pre-trained Transformer (ChatGPT) in developer-chatbot interactions within software development contexts. Using a combined dataset of over 10 000 conversations sourced from GitHub (commits, issues, discussions, pull requests, code files, repositories) and HackerNews, we apply text mining techniques like n-gram extraction, Sequential pattern mining (SPM), and topic modeling to analyze the functional expectations developers place on ChatGPT. The results reveal a strong contextual alignment between the source of the conversation and the role ChatGPT is assigned, ranging from scripting automation and debugging assistance to design planning and narrative generation. These findings demonstrate that developers use ChatGPT not merely as a code generator, but as a multifaceted collaborator.

Contents

1	Introduction	3
2	Background	5
2.1	ChatGPT and interaction with it	5
2.2	Related works	7
3	Data	9
3.1	Data cleaning	10
3.1.1	Code detection and its accuracy	10
3.1.2	Language detection	13
3.2	Data exploration	13
3.2.1	Number of prompts in conversations	14
3.2.2	Average symbol/word count per prompt	15
3.3	Data limitations	17
4	Methodology	18
4.1	Data preparation	18
4.2	N-grams	19
4.3	Sequential pattern mining	20
4.4	Topic modelling	20
5	Results	23
5.1	N-grams	23
5.2	Sequential pattern mining	25
5.3	Topic modelling	27
6	Discussion and conclusion	30
6.1	Future work	31
7	Reflection on research process	32
7.1	Research results	32
7.2	Process	32
7.3	Problems encountered	33
7.4	Takeaways	33

A	Appendix	41
A.1	N-grams WordClouds	41
A.2	Topic models	45
A.3	Topic models' WordClouds	49

Chapter 1

Introduction

Large Language Model (LLM) is a statistical model that is trained on the large amount of data containing billions of parameters and that uses previous tokens to predict future tokens [1]. LLMs are used in different domains such as conversation agents, education, text summarization, explainable AI, information retrieval and many more [2]. One of the most well-known models is Chat Generative Pre-trained Transformer (ChatGPT) [3] — a chatbot developed by OpenAI and first released in November 2022. According to OpenAI CEO Sam Altman, ChatGPT reached 100 million weekly active users in November 2023¹.

ChatGPT drew attention for its detailed and expressive answers across many domains of knowledge, including software development. The chatbot was adopted by software engineers as an alternative to using Google for solving software development-related issues. Studies have investigated ChatGPT’s efficacy when used in software development processes [4], its limitations [1], domains of application within issue-tracking systems, reliance on the ChatGPT-generated code [5] and the correctness of the generated code [6, 7]. Additionally, some studies focus on human-bot interaction: the role of ChatGPT in the process of collaborative software architecting [8], or what ChatGPT’s primary use is in issue-tracking systems [5].

In this study we focus on the domain of human-bot interaction and use developer prompts to extract the types of collaboration that developers want ChatGPT to follow during their interaction. There are many ways developers can collaborate with ChatGPT, and we focus on extracting patterns from their prompts to recognize what type of collaboration the developer expects from ChatGPT. This leads to the research question of this study: **What role does ChatGPT take in user-chatbot interaction within the software development domain?**

To answer the research question, we used various supervised and un-

¹<https://www.theverge.com/2023/11/6/23948619/openai-chatgpt-devday-developer-conference-news>

supervised text mining techniques: n-grams extraction, Sequential pattern mining (SPM) and topic modeling. These methods help us with understanding the common themes and topics present in developers' ChatGPT prompts and the expectations they place on ChatGPT in these prompts. Topic modeling provides information that describes the data, while SPM and n-grams are used to provide the initial overview of what are the most common words/word combinations present in the data. We apply the mentioned techniques on DevGPT [9] and its extension datasets that consist of developer-ChatGPT conversations sourced from GitHub² commits, pull requests, code files, issues and discussion, and from HackerNews³.

The answer to this research question provides valuable insights into the role that ChatGPT plays in developer-chatbot collaboration. Additionally, it shows what tasks developers trust chatbots like ChatGPT to perform and what to expect from ChatGPT in that role. By outlining the patterns of developer-chatbot interaction, we provide the ground for new research to develop or enhance chatbots' ability to help the software developer in the way that is most expected from the chatbot.

We firstly provide some background information and related research in Chapter 2.2. Then we describe in detail what the dataset consists of and what pre-processing steps are taken in Chapter 3. Chapter 4 focuses on methods used to explore the data. Chapter 5 contains the results of the research and their analysis. Chapter 6 concludes the research by discussing the methods and results of the study, and establishing the future research.

²<https://github.com>

³<https://news.ycombinator.com>

Chapter 2

Background

2.1 ChatGPT and interaction with it

Chat Generative Pre-trained Transformer (ChatGPT) is a large-scale, pre-trained language model developed by OpenAI and designed to engage in natural language conversations, utilizing Natural Language Processing (NLP) techniques that enable the generation of human-like text with context and coherence. Its performance in generating creative and novel text and ability to adapt to the context opened the possibilities for applications in various fields, such as creative writing, marketing, advertising, technical problem-solving, etc [10].

ChatGPT model is a transformer-based neural network [11] crafted to accomplish natural language processing tasks. Initially, the prototype was launched on the 30th of November 2022, which became available to the public on the 30th of January 2023. The model is trained on a vast corpus of textual data that was gathered from various sources, such as articles, books, reviews, online conversations and other human-generated data. This allowed the model to capture natural language patterns, nuances, and complexities, allowing ChatGPT to perform well in non-trivial dialogues on various topics.

This research focuses on the interactions with ChatGPT. Interaction with ChatGPT consists of *conversation*, which is a collection of pairs of user *prompt* and ChatGPT's reply. The conversation starts with the user creating the initial *prompt* which usually consists of the problem description, question, background information, step-by-step instructions for ChatGPT or any other information relevant to the user's needs. ChatGPT uses the information provided in the prompt to generate the *response*. After the user receives the response, they can either stop the conversation or continue it by providing the new prompt for ChatGPT to process. Such ChatGPT conversations are used in this research to answer the posed research questions with user prompt analysis.

Figure 2.1 shows the example of a developer-ChatGPT conversation.

The dark-blue square outlines the prompt of the conversation. The dark-green outline shows the full response of ChatGPT, which contains a code example outlined with orange, an explanation in the yellow outline, and a comparison to alternatives in the light-blue outline. The conversation can continue with the next prompt typed in the text input shown by the red outline.



Figure 2.1: Example of a ChatGPT conversation.

2.2 Related works

DevGPT [9] — the dataset used for this research — is a well-researched dataset that provides various insights into the developer-ChatGPT interactions and the usability of ChatGPT in the field related to software engineering.

The DevGPT dataset provides great opportunities to evaluate how ChatGPT handles questions and requests related to code generation and debugging, and how well the generated code corresponds to the set standards of code security, cleanliness, and efficiency. Several studies analyzed various aspects of the generated code: some studies focused on the security aspect and many other studies on the quality aspect of the generated code. The studies done by Siddiq, et al. [12] and Rabbi, et al. [13] revealed that the generated and modified code contained several CWE (Common Weakness Enumeration)¹ vulnerabilities, and issues captured by the Pylint 3.0.2² code quality analysis tool. Moreover, Zhang, et al. [14] showed that at least 35% of the 98 generated Kubernetes manifests contained at least one instance of a code smell. Other studies that researched the generated Python and Java code have discovered various code safety and quality issues by analyzing the generated code [15, 16, 17]. Moratis et al. [16] and Siddiq et al. [17] investigated the direct use of generated code in the source code base and concluded that ChatGPT generated code has to be reviewed and often modified if it is to comply with the code safety and quality standards.

Furthermore, code generation and code refactoring tasks have been addressed in more studies based on the DevGPT dataset. Grewal, et al. [18] and Jin, et al. [19] investigated its use for code generation, concluding that while ChatGPT can produce functional code, developers often need to verify its correctness and adjust the structure to meet the functional requirements. Code refactoring, on the other hand, is a more open-ended task, where ChatGPT tends to focus on improving readability, maintainability, usability, and performance, as shown by Chavan, et al. [20] and AlOmar, et al. [21]. Their research revealed that developers typically adopt a straightforward strategy of copying the code and directly requesting a refactor, which resolves the issue in an average of 2.7 to 4.6 prompts. Prompt quality is a critical factor in successful interactions with ChatGPT [22, 23]. Mondal, et al. [22] found that vague or underspecified prompts, especially those requesting additional features without clear specifications, often lead to longer and less efficient conversations. Altogether, these studies suggest that while ChatGPT is a valuable tool for code-related tasks, its effectiveness heavily depends on prompt clarity and developer oversight.

Moreover, other studies have analyzed various aspects of usability and

¹<https://cwe.mitre.org>

²<https://pypi.org/project/pylint/3.0.2/>

integration of ChatGPT-generated code into multi-language software development and pull request workflows. Aguiar, et al. [24] found that over 75% of ChatGPT’s code suggestions used a different language than the host repository, with the highest mismatch in CSS (91.16%) and the lowest in C# (60.18%) and CodeQL (58.08%). Chouchen, et al. [25] examined the topics of ChatGPT-related pull request conversations, identifying common themes such as code generation, refactoring, bug fixing, concept explanation, DevOps, testing, and recommendations, and noted that ChatGPT was more frequently used in large, time-consuming pull requests. Watanabe, et al. [26] and Hao, et al. [27] found that developers generally responded positively to ChatGPT’s code review comments, with fewer than one-third of replies expressing disagreement. Additionally, Hao, et al. [28] observed that ChatGPT was used in various roles in GitHub pull requests and issues, including as an author, reviewer, and collaborator. Together, these studies highlight ChatGPT’s growing presence and generally positive reception in collaborative software development environments. However, the responses should be critically checked as they still tend to contain code snippets written in a non-native language.

Lastly, researchers explored the structure, topics, and outcomes of interactions between developers and ChatGPT. Several studies using the DevGPT dataset have focused on analyzing the developers’ intent and the subjects of their requests. The study conducted by Sagdic, et al. [29] have used BERTopic [30] topic modeling NLP technique to create topic clusters, which can be manually labeled and added to the topic’s taxonomy. They have discovered seventeen topics that can be placed in seven categories: advanced programming; front- and backend development; DevOps integration and practices; streaming, media and localization; data management; and other topics. A similar approach was used by Mohamed, et al. [31], where they used four different classifiers to categorize all conversations by their purposes based on the developer’s needs. Their research revealed 19 categories, out of which code generation and seeking general information accounted for almost a third (31.15%) of all the conversations. Champa, et al. [32] have used *facebook/bart-large-mnli* model for zero-shot classification [33] to categorize the types of tasks developers present to ChatGPT, where they discovered that quality management of Python code and commit issue resolution tasks were the categories developers sought the most assistance with. While all three studies share a common goal of categorizing developer interactions with ChatGPT, they differ in methodology and focus, ranging from topic modeling to classification based on intent, yet consistently highlight code generation and problem-solving as dominant use cases.

Chapter 3

Data

To answer the research question, we used the DevGPT dataset [9], which contains ChatGPT conversations shared on GitHub and Hacker News between late May 2023 and October 12, 2023. The dataset includes nine snapshots (latest: October 12, 2023), each comprising links to ChatGPT conversations found in GitHub commits, issues, discussions, pull requests, code files, and Hacker News articles. Considering the origins the conversations were sourced from, it is safe to assume that the prompts for these conversations were written by software developers.

In this study, we used data from all available snapshots. To ensure data quality, we removed duplicate entries based on shared ChatGPT conversation links. Additionally, we extended the DevGPT dataset with a new dataset collected in March 2025. This extension contains developer-chatbot conversations shared in the same sources but posted after the latest DevGPT snapshot. The structure of the extended dataset mirrors that of DevGPT to maintain consistency. The dataset was provided for this research by the author’s colleagues from the same research group.

Table 3.1 presents the number of unique conversations sourced from each data source, along with the remaining number of conversations after removing duplicates and non-English conversations.

	Commits	Issues	Discussions	Pull req.	Code files	Repositories	Hacker News	Total
DevGPT DR	670	516	59	268	2010	0	322	3845
Extension DR	881	1134	65	775	9536	10	805	13206
Combined DR	896	1235	113	813	9959	10	858	13884
Combined NE	861	990	89	672	7980	8	760	11360

Table 3.1: Number of conversations collected from different sources after all the duplicates (DR) and non-English conversations (NE) were removed from the data.

The total number of conversations shared across all the sources after duplicates and non-English conversations were removed is 11360. However,

after removing duplicate conversations from the entire combined dataset, it leaves us with 10175 conversations, which means that 1185 developer-ChatGPT conversations are shared in multiple locations.

3.1 Data cleaning

During our initial analysis of the dataset, we identified the presence of information irrelevant to the research question — referred to as “noise”. First, many prompts contain embedded code alongside natural language. A good example of this is the following prompt from the DevGPT dataset sourced from GitHub Issues:

So you want me to hard code the password in this section of src/pages/api/getDeviceInfo.js?

```
const pool = new Pool({
  host : process.env.PGHOST_2,
  port : process.env.PGPORT_2,
  database : process.env.PGDATABASE_2,
  user : process.env.PGUSER_2,
});
```

Second, while we focus exclusively on English-language conversations, the dataset includes interactions in multiple languages. For example, the following prompt was sourced from conversation shared in GitHub commits: *mach mir noch eine aufgabe zu dem selben muster, damit ich das thema besser verstehen kann.*

Both of the mentioned factors introduce noise that must be removed before proceeding with data exploration and answering the research questions.

3.1.1 Code detection and its accuracy

Initially, we attempted to detect programming language content using the Guesslang Python library¹, but its accuracy and F1-score were too low for our dataset. As a result, we developed a custom method for identifying and filtering out programming lines.

To identify code segments within prompts, we first split each prompt into individual lines using the line break character (`\n`). We then applied a custom script to estimate the likelihood that each line contains code. This approach is based on the assumption that code blocks typically appear on separate lines from natural language text and often span multiple lines within a prompt. As a result, line-by-line analysis is necessary to accurately detect code.

¹<https://pypi.org/project/guesslang/>

To identify whether a line is written in a programming language or natural language, we implement a function that evaluates the line's structures, whether they are tokens or full sentences, and assigns them a probability score: 0 indicates a programming language structure, 0.5 indicates ambiguity (could be either), and 1 indicates natural language. The list below shows what kinds of structures are considered in this script:

- Line is empty (likelihood: 0);
- Line starts a multiline comment block with symbols like `/*`, `"""`, `'''`, or `<!--`, and does not end on the same line — the whole block is marked as comment (likelihood: 0);
- Line is within a multiline comment block and ends with symbols like `*/`, `"""`, `'''`, or `-->` — marked as comment (likelihood: 0);
- Line matches single-line comment patterns such as `//`, `#`, `--`, etc. — marked as comment (likelihood: 0);
- Line starts with an opening HTML tag (e.g. `<div>`) and does not end with a closing tag — line and subsequent lines until closing tag are marked as HTML (likelihood: 0);
- Line contains a complete HTML block (e.g. `<p>Text</p>`) — marked as HTML (likelihood: 0);
- Line is tokenized using `shlex.split` if it contains quotes, otherwise by whitespace. Each token (word) is assigned a probability score based on specific heuristics:
 - Word contains space and resembles a sentence — averaged likelihood from subword analysis;
 - Word starts with tab (likelihood: 0);
 - Word is in a list of known programming keywords (likelihood: 0.5);
 - Word is a camelCase or PascalCase identifier (likelihood: 0);
 - Word is in UPPERCASE (e.g. SQL keywords; likelihood: 0.5);
 - Word is purely alphabetical (likelihood: 1);
 - Word is a comment-like or inline programming pattern (e.g. `::`, `#`, `<\`, `<|`–; likelihood: 0);
 - Word matches function calls, method chains, array indexing, or regular programming syntax (e.g. `func(x)`, `obj.attr`, `arr[i]`; likelihood: 0);
 - Word ends with punctuation (e.g. `word,`, `word.`, `word:`; likelihood: 0.5);

- Word ends with semicolon often used in programming syntax (likelihood: 0);
- Word lacks alphabetic characters (likelihood: 0);
- All other cases (likelihood: 1).

To determine whether a line is written in natural language or code, the script calculates the average likelihood score by summing the individual scores of all structures (tokens or sentences) on the line and dividing by the total number of structures. If the resulting average is above 0.5, the line is considered to contain natural language; otherwise, it is classified as a code line.

The effectiveness of the code detection script was evaluated using ten randomly selected conversations, from which up to 20 lines were randomly sampled. These lines were processed by the script, and the results were saved to a file. Each line was then manually labeled as either code or non-code for comparison. The script demonstrated satisfactory performance across all tested conversation sources. Table 3.2 shows how well the code detection script works.

	Accuracy	F-score
GitHub commits	0.96	0.98
GitHub issues	0.92	0.95
GitHub discussions	0.92	0.93
GitHub pull requests	0.96	0.97
GitHub code files	0.95	0.96
GitHub repositories	0.91	0.94
HackerNews	0.92	0.94
Total	0.92	0.95

Table 3.2: Accuracy and F-score of the code detection script per data source category.

False-positives and false-negatives were examined to better understand the types of errors made by the code detection script. The most common false-negatives occurred with error messages or log outputs, which often contain large amounts of natural language, making them difficult to distinguish from regular text. On the contrary, sentences with a high number of punctuation marks were frequently misclassified as code, even when they were not. For instance, lines such as “*Done.*” and “*- Prefer async/await over promises!*” were incorrectly identified as code, although they could plausibly be part of natural language text.

This evaluation method also has limitations, particularly in handling multi-line HTML or comment blocks. Because lines for testing are selected randomly, it is possible to include a line that starts a multi-line HTML tag or comment block without capturing the corresponding closing line. As a result, the following lines may be incorrectly labeled as part of HTML or a comment. Additionally, some lines contain both code and natural language, making the results of the code detection depend on their proportion. Examples of such lines are questions that contain short code snippets in them. The presence of the code in such sentences might affect the research results and can be considered a limitation of the current data cleaning method.

Despite these limitations, the script achieved an average accuracy and F-score of over 90%, which we consider a good result for the purposes of this research.

3.1.2 Language detection

For language detection, which occurs after code lines were removed, we used two existing libraries: *langdetect*² and *pycld2* (*Compact Language Detect 2*)³. Using both libraries improves detection accuracy: if one library fails to identify the language, the result from the other helps determine whether the prompt should be kept (if detected as English) or discarded.

The language identification process operates at the line level: each prompt is split into individual lines, and the language is identified for each line separately. If a line is detected as containing only programming code (as described in Section 3.1.1), it is marked as code and excluded from language detection. This approach helps avoid misclassifying non-English prompts as English due to the presence of English keywords commonly found in code.

The output of the language check is a tuple, where the first element represents the result from *langdetect*, and the second from *pycld2*. If the most common tuple across all lines in the prompt is either (“en”, “en”), (“en”, “un”), or (“un”, “en”), the prompt is considered to be in English. Any other combination results in the prompt being discarded.

3.2 Data exploration

As part of the data exploration process, we begin by extracting key information about the dataset. For the purposes of this study, we focus exclusively on the user-generated prompts and disregard the responses produced by ChatGPT across all conversations.

²<https://pypi.org/project/langdetect/>

³<https://pypi.org/project/pycld2/>

3.2.1 Number of prompts in conversations

To better understand the structure and variability of the dataset, we examine the distribution of conversation lengths across different data sources. Table 3.3 shows the basic statistics of conversation lengths across different data sources. Most datasets contain short conversations on average, with median lengths between 1 and 2 prompts, suggesting that interactions are typically brief. However, GitHub repositories stand out with a higher mean (21.38) and median (10.5), as well as the largest standard deviation (31.58), indicating a wider spread and the presence of longer discussions. This phenomenon can be explained by a low number of sharings (8 after non-English conversations’ removal), where each heavily influences the data.

Source	Mean	Median	Standard deviation
GitHub Commits	3.2	1	5.65
GitHub Issues	4.57	2	8.9
GitHub Discussions	4.66	2	9.52
GitHub Pull requests	3.36	2	7.94
GitHub code files	6.72	2	16.74
GitHub repositories	21.38	10.5	31.58
HackerNews	4	1	9.12

Table 3.3: Median, mean and standard deviation of the conversation lengths from different sources.

Figure 3.1 contains histograms that show how many conversations contain a certain number of prompts, where the conversations with the maximum number of 30 prompts are shown. All the conversations that exceeded the number of 30 prompts were separated into separate JSON files and their content was investigated and described below.




Figure 3.1: Lengths of conversations in histogram gathered from different sources.

Long conversations contents

To understand the reasons behind unusually long conversations, we manually reviewed 295 conversations that exceeded a predefined length threshold. The most common patterns are summarized below:

- **Pasted content:** Users split large texts (e.g., code files, articles, books) into multiple prompts due to the prompt’s character limits.
- **Iterative debugging:** Users refine requests or see help with errors after initial code responses did not meet their expectations.
- **Programming assistance:** Users guide ChatGPT through multi-step coding tasks, testing and iterating on output.
- **Mentorship:** Users share their own code for feedback, explanations, or improvement suggestions.
- **General discussion:** Conversations focus on non-programming topics such as advice, opinions, or general discussions.

3.2.2 Average symbol/word count per prompt

The average number of symbols per prompt varied across sources, as shown in Figure 3.2. The x -axis represents the prompt position within a conversation, while the y -axis shows the average symbol count for that position. Since longer conversations are less common, later prompt positions often reflect fewer data points, leading to the prompt average pattern changes that can be visibly observed in the conversations sourced from GitHub code

files or Hacker News. Across the majority of the sources, initial prompts typically have higher symbol counts, as they often introduce the problem in detail.




Figure 3.2: Average symbol count per prompt number for all conversation sources after code lines were removed.

To identify outliers, we defined dataset-specific cut-off thresholds set at two-thirds of the maximum observed average. Conversations containing prompts exceeding these thresholds were saved for further analysis.

Long prompt contents

Conversations containing unusually long prompts were saved and manually reviewed to understand their content. In total, 1984 conversations contained long prompts. We identified several common reasons for the prompt length:

- **Pasted external content:** Users copy large texts from sources such as scientific articles, GitHub issues, or documentation, often asking for summaries, explanations, or audience-specific rewrites.
- **Code-related debugging:** Prompts include error messages, undocumented code, or comments not flagged as code, with users seeking help with debugging or improvement.
- **Detailed problem descriptions:** Some users provide extensive background to frame their query, including:
 - Highly specific requests or detailed issue descriptions;
 - Roleplay-style prompts (e.g., asking Chat Generative Pre-trained Transformer (ChatGPT) to act as a mentor or expert);

- Contextual or background information;
- Multiple follow-up questions or clarifications;
- Previously attempted (but unsuccessful) solutions;
- Long URLs or resource links;
- Lists of requirements or instructions;
- Desired output examples or target functionality.

3.3 Data limitations

The dataset used for this study offers a useful overview of how developers utilize ChatGPT to address everyday challenges. However, the data sources are incomplete. Conversations were collected from GitHub (commits, code files, discussions, pull requests and issues) and HackerNews, covering conversations that happened before March 2025. As a result, the dataset leaves out other platforms that may contain additional insights into developers’ use of ChatGPT. Moreover, the data spans nearly two years, and some older conversation links that were used for scraping the data are now inactive, making it impossible to verify the original content.

The language detection and data cleaning tools used were also imperfect. They occasionally failed to identify prompts or prompt lines containing foreign languages or embedded code. Many prompts contain pasted content (articles, documentation, code comments, or error messages), which, although written in the natural English language, lacked relevance to the research questions. These elements introduced noise that the current toolset could not fully filter out.

Chapter 4

Methodology

The purpose of investigating how developers use ChatGPT in their projects is to gain a deeper understanding of the role the chatbot plays in developer-chatbot collaboration. This understanding can help improve the chatbot’s ability to assist developers more effectively. To extract this knowledge, we applied n-grams, Sequential pattern mining (SPM), and topic modeling techniques, which are described in this chapter.

4.1 Data preparation

To enable further analysis, we pre-processed the cleaned data by lemmatizing and tokenizing each prompt.

The lemmatization process involves splitting each sentence into a list of elements such as words, punctuation marks, digits, or alphanumeric strings, and then converting each word to its lemma using the *lemmatize_sentence* function from the **pywsd** package [34]. The function attempts to convert each word into its base form. If the lemma is not found in WordNet, the function applies stemming as an alternative processing step.

Once all the sentences are lemmatized, they are cleaned from the non-relevant data. We remove the following tokens:

- **Non-ASCII tokens;**
- **Hyperlinks** — token starts with “*http*”;
- **Stop words** — based on the `stop_words` list from the **nlTK** package¹ [35];
- **Punctuation marks;**
- **Tokens containing digits.**

¹Stop words list

The resulting lemmatized and cleaned lists of tokens are then used for the following tasks, including n-gram extraction, sequence mining, and topic modeling.

4.2 N-grams

N-grams is a sequence of n consecutive characters, words or other textual structures [36], where n refers to the number of elements in each group. These sequences are generated by sliding a window of size n over the text, preserving the original order of the elements. In this research, we focus on word-based n-grams. For example, the sentence “*That fox is chasing a rabbit*” can be tokenized into the following sequence of words: $\{that, fox, is, chasing, a, rabbit\}$. From this sequence, the following n-grams can be extracted:

- Bi-grams: $(that, fox), (fox, is), (is, chasing), (chasing, a), (a, rabbit)$;
- Tri-grams: $(that, fox, is), (fox, is, chasing), (is, chasing, a), (chasing, a, rabbit)$;
- Quadri-grams: $(that, fox, is, chasing), (fox, is, chasing, a), (is, chasing, a, rabbit)$;

Before the n-gram extraction, the data is preprocessed as described in Section 4.1. This includes the removal of stop words — commonly used words that carry little semantic content “*that*”, “*is*” and “*a*”. In the example above, these words would be excluded, leaving only content-bearing words for n-gram construction.

N-grams are extracted using the `ngrams()` function from the **nlTK** package, with the parameter n set to the desirable value. The extracted n-grams are then used to compute a frequency distribution, allowing us to sort them in descending order of occurrence. Such ranking helps highlight the most common lexical patterns found in the dataset.

For easier visualization, we use the WordCloud package [37] to generate visual representations of the n-gram frequencies. In these visualizations, the font size of each n-gram corresponds to its relative frequency where more frequent patterns appear larger.

During the initial extraction process, we observed that some tri-grams and quadri-grams contained repeated instances of the same word (e.g., $(position, position, position)$). This effect is likely caused by the combination of lemmatization/stemming, stop word removal and imperfect code line detection script, which can cause sequences of otherwise distinct words to collapse into repeated lemmas. To improve our results, we chose to exclude any n-grams that contain the same token repeated three or more times. This filtering step helps the quality of the analysis while preserving the patterns in the dataset.

4.3 Sequential pattern mining

Sequential pattern mining (SPM) is an essential part of knowledge discovery and data analysis often used in the analysis of medical treatment history [38], customer purchases [39] and digital clickstream [40]. SPM is used to find frequent patterns in large sequence databases.

An input for SPM is a set of sequences referred to as *sequence database*. Each sequence is an ordered set of *items*. A *pattern* is an ordered sub-sequence that occurs in at least one sequence in the database. *Frequency* defines the number of sequences that contain a pattern. Given an input sequence database, SPM involves finding frequent patterns that occur more than the specified threshold [41].

For example, if we have the sequences $\langle A, A, B, A, D \rangle$, $\langle C, B, A \rangle$ and $\langle C, A, C, D \rangle$, and the minimum frequency is set to two, the SPM extracts the following sequences: $\langle A, D \rangle$, $\langle B, A \rangle$ and $\langle C, A \rangle$.

In this research, we use **Seq2Pat** package [42] in order to extend sequential pattern analysis. N-grams, discussed in Section 4.2, extract sequential data of the set size. However, we expand this analysis using SPM in order to analyze frequent patterns seen in data that are longer than the longest n-grams we extract.

In order to extract frequent patterns, we initialize the *Seq2Pat* object with tokenized prompts as a sequence database, set maximum span to 10, batch size to 10000 and discount factor to 0.2. We extract patterns using the *get_patterns()* function, where the minimum frequency is set to the length of the database divided by 50, meaning that 2 or more percent of the data contains the sequence. Once the sequence patterns are extracted, they are sorted in descending order and saved for future analysis. The minimum frequencies for each data source were the following:

- GitHub commits: 51;
- GitHub issues: 78;
- GitHub discussions: 8;
- GitHub pull requests: 42;
- GitHub code files: 981;
- GitHub repositories: 3;
- Hacker News: 52.

4.4 Topic modelling

Topic models are generative models that provide a probabilistic framework for the term frequency occurrences in documents in a given corpus [43].

These models are used for large electronic text collections to organize, understand, search for, and summarize their content.

The topics are the relations that link words in a vocabulary and their occurrence in documents. Each document is viewed as a mixture of topics. Topic models discover different themes present in the text collections and annotate the documents according to the themes. The document coverage distribution of topics is generated to provide an overview of topics found in this document collection.

In this study we used Latent Dirichlet Allocation (LDA) topic modeling — one of the most popular topic modeling techniques. LDA is a probabilistic generative model that allows the observations to be described by unobserved data that explains why some parts of the data are similar to each other [44]. In LDA, documents consist of multiple topics and each topic is a distribution over a fixed vocabulary. For example, if we have the following vocabulary of $\{pan, cook, football, basketball\}$, then the kitchen topic will assign high probabilities to the words *pan* and *cook*, while *football* and *basketball* will have very low probabilities; however, the sport topic will have the opposite probabilities for all the words.

To determine the optimal number of topics for each data source, we used both coherence [45] and perplexity [46] scores. The coherence score measures the semantic similarity, with values ranging between 0 and 1, between the top words in each topic, reflecting the interpretability and quality of the topics. The perplexity score evaluates how well the model predicts unseen data, with lower values indicating better generalization. We trained topic models with 2 to 10 topics and computed both scores for each model. The number of topics that achieved the highest coherence score and lowest perplexity was considered to best represent the data and was selected for further research. Table 4.1 shows the number of topics selected for the data from each source.

Source	Number of topics	Coherence	Perplexity
GitHub Commits	4	0.4157	16386
GitHub Issues	2	0.4341	183417
GitHub Discussions	9	0.1844	44488
GitHub Pull requests	10	0.401	86206
GitHub code files	5	0.4367	2847
GitHub repositories	2	0.4321	6255
HackerNews	8	0.2758	104586

Table 4.1: Selected number of topics for data from each data source with coherence and perplexity scores.

For the data from most data sources, the highest found coherence score

was slightly above 0.4. However, topic models for the data sourced from GitHub discussions and HackerNews had the highest perplexity of 0.18 and 0.28 respectively. Such a low coherence score indicates that the quality of topics is poor. Additionally, the perplexity of the topic models for the data sourced from GitHub issues, pull requests and HackerNews shows quite high numbers, meaning that the topic models might not describe the unseen data well enough. Thus, topic models for conversations sourced from GitHub issues, discussions and HackerNews should be considered with caution.

In the implementation we use the data cleaned in the pre-processing step as described in Section 4.1. We start with creating the vocabulary: this is done using the *CountVectorizer* class to convert a collection of text documents to a matrix of token counts [47]. The parameters of *min_df* and *max_df* are set to 0 and 1 respectively; *ngram_range* is set to tuple of (1,4), making vocabulary to consist of n-grams of the length 1 to 4 words; and the stop words are the stop words from **nlTK** package. For topic modeling, we use the *LatentDirichletAllocation* class from **sklearn** package [48], where we set the number of topics to the number with the highest coherence and low perplexity, iteration count to 10 and random state to 42. We calculate the perplexity and coherence score of the topic model and visualize them along with WordCloud visualization of the topics.

Chapter 5

Results

In this chapter we provide insights into the knowledge we managed to extract during the analysis of developer’s prompts in developer-ChatGPT interaction. This chapter focuses on the results extracted from N-grams, topic modeling and sequence mining of the prompts.

5.1 N-grams

N-grams are extracted in order to provide an overview of the most common lexical patterns found within the data. Data gathered from each source was analyzed separately and the most common 1- to 4-grams were extracted. Table 5.1 shows interesting tri- and quadri-grams extracted from various data sources.

Source	n	Found n-grams
GitHub commits	3	shell script create, encode enclose result, avoid use sed, implement following feature
	4	script create change file, avoid use sed favor, task implement following feature
GitHub issues	3	th block th, valid string generation, state lookahead token, kernel patch instal
	4	td th block th, string generation metadata information, metadata information visit input, community edition line module
GitHub discussions	3	technology feasibility stable, operational feasibility challenge, clause potentially unfair, employee desirability mixed/high, open/close code environment

	4	technology feasibility stable employee, employee desirability high team
GitHub pull requests	3	gazebo sensor imu, plugin sensor accelerometer, deprecate pip enforce, possible replacement use
	4	coelectric cavity lumen subclassof, gazebo sensor imu plugin, sensor accelerometer parent origin, replacement use pip package, string sure want delete
GitHub code files	3	rdf type owl, restriction owl minqualifiedcardinality, passage book rewrite, break fourth wall
	4	rdf type owl objectproperty, rdfs subpropertyof owl topobjectproperty, george r r martin, book rewrite passage remove, break fourth wall include, wall include reference reader
GitHub repositories	3	fill height viewport, increase understanding universe, help reduce suffering, research educational institution
	4	height fill height viewport, leave position foreground scyscraper, reduce suffering increase prosperity, prosperity increase understanding universe, planet intellectual community train
Hacker News	3	capture position position, form line capture, tic tac toe, session user host, level level headline
	4	form line capture position, x currently form line, headline insert concat star, concat star schedule progn

Table 5.1: Interesting 3- and 4-grams extracted from the data.

The extracted n-grams reveal distinct linguistic and thematic patterns across different data sources. Prompts from GitHub commits and pull requests use task-oriented, direct and action-driven communication typical of development workflows. The prompts often focus on modifying, automating or improving code artefacts. In contrast, GitHub discussions contain planning-related n-grams like “clause potentially unfair” and “technology feasibility stable employee”, suggesting that contributors debate the feasibility and desirability of the solution, not just technical correctness.

GitHub repositories show recurring themes in UI rendering (e.g.m CSS and Viewport references) as well as humanitarian, educational, and scientific discussions. However, prompts from GitHub issues and code files are influenced by embedded code or markup, such as HTML and RDF/OWL artefacts, indicating that developers often seek help with specific technical

problems.

Hacker News conversations differ in a less technical, more playful approach. N-grams like “tic tac toe”, “form line capture position”, and “headline insert concat star” combine elements of game logic, syntax play, and UI behavior, typical of exploratory or playful community posts.

In conclusion, each platform reflects its users’ intent, ranging from practical coding to creative exploration.

5.2 Sequential pattern mining

Sequential pattern mining (SPM) is used to extend sequential pattern analysis of the data present in the datasets. For this research we consider sequences that are longer than 4 tokens (1-, 2-, 3- and 4-grams were analyzed in Section 5.1) and appear in 2 or more percent of the non-empty prompts.

These requirements were applied to the data sourced from all 7 sources. However, only data sourced from GitHub commits and repositories contained sequences that corresponded to the requirements. The data sourced from GitHub commits, issues, code files, pull requests and HackerNews contained no such sequences. Thus, in this part of the research we only focus on the data from the two sources and their contents.

Table 5.2 shows categories with the most informative and recurrent sequence patterns extracted from GitHub commits and repository content. The results are grouped by source and functional themes.

Source	Group	Found sequences
GitHub commits	Top-5	format shell script create everything (660), format shell script create change file everything (656), output format shell script create change file (656), format script create change file everything (656), format shell create change file everything (656)
	Scripting automation	enclose result shell script create file solve task (645), format encode enclose result shell script create change file everything (645), shell script create change file everything solve task file small (605), export single output format encode result shell script create (121)
	Constraints and guidance	avoid use sed favor full file use prevent substitution write (250), avoid use sed full file solution write concise (119)

	Feature implementation	implement feature create new file need file (78), implement following feature create file need file (78), project specific file export single function (78)
	Platform dependencies	task file small full file osx (78), substitution osx npm jq write text outside script (77)
GitHub repositories	Top-5	fill height viewport leave position foreground left edge right (10), height fill height viewport leave position foreground left edge (10), height viewport leave position foreground left edge right position (10), viewport leave position foreground left edge right position foreground (10), fill height leave position foreground left edge right (10)
	Web UI layout composition	fill height viewport leave position foreground left edge right (10), fill height viewport leave position foreground skyscraper left edge right (9), foreground skyscraper left edge right position foreground skyscraper right edge (9), landscape foundation height set height fill viewport skyscraper height fill (9)
	Continuous research benefits	continue research work help reduce suffering planet intellectual community restore (8), disaster preparedness program action help reduce suffering increase prosperity increase (8), planet share knowledge expertise planet scientist scholar help fill knowledge (8)
	DOM structure	stylesheet body div landscape div sun div city skyscraper div (3), print data get document embed id metadata data document data (3), embed id metadata document embed convert list (3)

Table 5.2: Interesting sequences found in the data sourced from GitHub commits and repositories.

The sequence data from conversations sourced from GitHub commits and repositories shows a separation in focus: commits envode more procedure-related requests, while repositories focus more on structural and layout information.

The most frequent patterns in GitHub commit-related conversations focus on shell scripting for file manipulation and task execution, and its constraints. Repetitive sequences like *“enclose result shell script create file*

solve task” and *“format encode enclose result shell script create change file everything”* indicate a multistep automation pipeline for file manipulation, including file formatting, encoding, and task resolution. Additionally, commit data provides constraints aimed at maintainability, with sequences discouraging tools like **sed** in favor of full-file edits. Other patterns describe feature development practices, such as modular file creation, and even include platform-specific workflows using tools like **npm**, **jq**, or referencing **osx**.

In contrast, conversations sourced from GitHub repositories contained sequences that are dominated by UI layout and rendering logic. Patterns like *“fill height viewport leave position foreground left edge right”* reflect deterministic spatial composition, likely connected to DOM or canvas rendering engines. References to *“skyscraper”* — a vertical stacking model, and patterns containing explicit DOM hierarchies, reinforce a component-based layout approach.

Overall, conversations from commits capture the execution trace of development, like scripts, tasks, and automation, while conversations from repositories reflect the architectural blueprint of the system, including layout logic and project narrative.

5.3 Topic modelling

To investigate how developers interact with ChatGPT across different development contexts, we used topic modeling on conversations extracted from seven sources, including GitHub-based sources (commits, issues, discussions, pull requests, code files, and repositories) and HackerNews threads. Latent Dirichlet Allocation (LDA) topic modeling was used in order to analyze the data and extract topics covered in the data. The optimal number of topics for each part of the dataset was determined by selecting the number of topics with high coherence and low perplexity scores. Section 4.4 covered this process in more detail.

Appendices A.2 and A.3 provide a full overview of the topic models in table and WordCloud form, respectively. The resulting topic models, summarized in Tables A.1 to A.7, reveal a distribution of themes aligned with different stages of the software development lifecycle.

The topics derived from commit-based conversations are focused on actionable programming tasks, like file handling and project configuration. The **Task-based file management** topic includes words like *“file”*, *“task”*, *“solve”*, and *“create”*, which point to developer requests aimed at automating or scripting low-level actions. Similarly, **Function testing** and **UI layout planning** reflect frequent use cases in which developers ask ChatGPT for code to validate behavior or generate visual components. The **AI-assisted development** topic suggests that developers outsource code scaffolding.

folding to ChatGPT, especially for feature generation and scripting tasks.

In issue-oriented conversations, developers focus more on troubleshooting behavior and reasoning about data. Topics such as **Code & model behavior** and **Data & value processing** contain terms like “*error*”, “*input*”, “*model*”, and “*value*”, showing that developers use ChatGPT for clarification and debugging assistance when encountering code errors or inconsistencies. These topics underscore ChatGPT’s use as a debugging companion and computational reasoning assistant.

Discussions are broader in scope and, as noted before, have topics with a low coherence score. The **Product design feasibility** topic includes terms such as “*feasibility*”, “*desirability*”, and “*technology*”, indicating higher-level conversations about user needs or system requirements. Several topics contain references to the *Silicon Valley* series¹, mentioning characters and companies’ names. Technical topics like **Node installation issues** and **Model calibration** show that developers also engage ChatGPT in discussions about environment setup and machine learning workflows, respectively. These themes blend the topics of software engineering, storytelling, and machine learning in discussion forums.

Conversations sourced from pull requests are described with technical precision, and deal heavily with code correctness, dependency resolution, and language semantics. Topics such as **C++ function signatures**, **Ontology & classification**, and **Mathematical constants & types** suggest that developers use ChatGPT to verify and generate low-level code. Other topics like **Web attestation APIs**, **Node package metadata**, and **Python packaging errors** reflect common questions about system integration, library usage, and environment misconfiguration. These patterns show how developers rely on ChatGPT for help in aligning implementation details with expected behavior.

Code file-based conversations reflect on documentation, algorithmic generation, and knowledge encoding. The **Step-by-step coding guidance** topic (e.g., “*step*”, “*please*”, “*solution*”) shows that many conversations include instructional sequences. The presence of **OWL/RDF ontologies** suggests ChatGPT use in data modeling, and **Functional file operations** shows traditional programming tasks assistance. Meanwhile, **Evidence-supported argument** and **Book passage creation** indicate mixed-domain usage, showing that ChatGPT also supports non-code-related writing tasks within GitHub codebases.

Topics from repositories’ conversations focus on structured layouts and project-level semantics. The **UI layout metadata** topic includes tokens such as “*viewport*”, “*skyscraper*”, and “*fill*”, reflecting UI generation and interface rendering. A second topic, **Research effort**, contains mission-driven language (e.g., “*planet*”, “*disaster*”, “*reduce*”), indicating that repository-

¹<https://www.imdb.com/title/tt2575988/>

level content includes vision statements or long-term goals.

Topics extracted from HackerNews conversations reveal a more dialogic tone compared to other platforms. **Conversational requests** and **Narrative creation** show how developers use ChatGPT to simulate discussions and tell stories. Others, like **Game interaction errors**, **File tracking utilities** and **Code support**, return to utility-based interactions. Topics like **Company schedule** and **Positional pattern capture** contain many repetitive sequences, making them uninformative. These sourced conversations appear to blend technical querying with user-driven storytelling, positioning ChatGPT as both an assistant and a conversational partner in exploratory thinking.

Across all seven data sources, topic modeling reveals that developer-ChatGPT interactions are highly context-sensitive, aligning closely with the purpose and communication norms of each platform. While conversations sourced from GitHub commits and issues emphasize low-level execution and debugging, conversations sourced from pull requests and code files highlight correctness and structure. In contrast, ChatGPT conversations found in GitHub discussions and repositories tend to encode project-level reasoning and documentation, and HackerNews reflects informal, exploratory, or social discussions. Collectively, these patterns underscore ChatGPT’s role as a multifunctional collaborator: serving as a code assistant, design partner, documentation helper, and narrative generator.

Chapter 6

Discussion and conclusion

This study investigated the roles that ChatGPT plays in conversations with software developers, using a large, multi-source dataset of developer-ChatGPT interactions. Using a combination of n-gram analysis, Sequential pattern mining (SPM), and topic modeling across seven distinct data sources, which include GitHub commits, issues, discussions, pull requests, code files, repositories, and HackerNews, we studied the lexical and thematic structures that characterize developer expectations and chatbot functionalities.

Key findings show contextual dependence of ChatGPT roles, derived by the origin of the conversation. For example, GitHub commits mainly reflect task automation and file-based scripting workflows, which supports the view that developers treat ChatGPT as a procedural collaborator, trusted with command-line logic, automated code generation, and low-level implementation tasks. These findings are in-line with the work of Sagdic et al. [29] and Mohamed et al. [31], who also observed dominant use cases in code generation and automation.

Contrarily, conversations sourced from GitHub issues and discussions show a different interaction style. Topics in these sources express debugging logic, system configuration, and architectural reasoning. Discussion threads reflected higher-level planning and storytelling, with themes like **Product design feasibility** and **Silicon Valley series** showing that ChatGPT is sometimes treated as a reasoning or storytelling partner, and not only as a programming assistant.

Conversations from GitHub pull requests and code files illustrated the importance of semantic and syntactic correctness. Developers often use ChatGPT for tasks involving function validation, API usage, and type system clarification. These contexts demand higher trust in the chatbot’s technical accuracy, yet the literature cautions against full automation due to persisting issues with correctness and code safety [12, 13, 15].

Finally, conversations sourced from GitHub repositories and HackerNews

revealed ChatGPT’s capacity to support documentation, vision statements, and informal discourse. Sequences like *“fill height viewport leave position foreground left edge right”* suggest that ChatGPT is used to assist in UI layout composition and document structuring. Meanwhile, HackerNews conversations uncovered more dialogic and playful use cases, showing ChatGPT serving as a narrative generator, ideation tool, and storytelling agent. This supports previous research emphasizing the model’s capabilities in non-code domains [10].

From a methodological perspective, the integration of SPM and topic modeling yielded some interpretable patterns. The findings demonstrate how varied ChatGPT’s role is in software development, ranging across domains from automation and testing to planning and narrative creation. These findings contribute to the research question posed in this study by showing that ChatGPT’s perceived role spans over code assistant, design partner, documentation helper, and conversational companion.

6.1 Future work

Future research could extend this analysis by incorporating ChatGPT responses, allowing for a more holistic understanding of co-construction and feedback loops in developer-chatbot dialogues. Additionally, fine-grained intent classification and temporal analysis could reveal evolving usage trends and adaptation behaviors. Lastly, there is potential to design tailored Large Language Models or prompt-engineering guidelines optimized for the distinct roles developers expect ChatGPT or other chatbots to fulfill.

Chapter 7

Reflection on research process

7.1 Research results

In conclusion, the chosen methodologies were effective in producing interpretable results from the collected developer-ChatGPT conversations. All approaches returned core lexican and sequential patterns found in the data, and the similarity of results across methods strengthened the confidence in their validity. Although the applied methods rely heavily on frequent words and phrases, the use of multiple methods enabled verification of findings from different angles.

While more advanced approaches such as BERTopic or other transformer-based methods could have found more subtle or abstract patterns, these were not applied due to limited knowledge in deep learning. Instead, I prioritized transparency and control through well-understood, classical techniques. Future work may benefit from incorporating neural models to explore more subtle and less frequent interaction patterns.

7.2 Process

The research internship began in February 2024 with the selection of the topic in collaboration with the supervisor. The research methods were defined early in the process, followed by the implementation phase, which occupied the first three months. During this period, the main components were developed and refined, and the initial drafting of the accompanying research article was initiated. After this initial phase, the project went through an extended pause until April 2025. When the project restarted, the focus was on finalizing the implementation and completing the research article.

7.3 Problems encountered

There were no significant technical or methodological difficulties encountered during the research process. The selected approaches allowed the generation of coherent and interpretable results. Both the implementation of analytical methods and the interpretation and documentation of findings did not cause any issues.

However, the primary limitation affecting the progress of the research was external: personal life circumstances, followed by a master's thesis, led to a need to pause the internship until the thesis was done. As a result, the overall timeline of the project was delayed. This pause impacted momentum and reduced the originality of the research, as several new publications released during that period that addressed related themes in developer-chatbot interaction and the role of ChatGPT in these interactions.

7.4 Takeaways

One of the main takeaways from the internship experience was the importance of timing and planning. Firstly, the research plans were created for both the internship and the master's thesis, but the experience gained during the internship influenced how I approached the thesis. The delays I encountered early in the internship, mainly due to not fully following the initial plan, helped me recognize the risks of losing momentum. As a result, I was able to follow the thesis planning better and manage the project more efficiently. Secondly, I have used the knowledge of the value of clearly defining the research question and selecting appropriate methodologies early in the process. This helped me focus on the most relevant parts of the work and ensure that both the implementation and analysis were aligned with the project goals.

Moreover, my master's thesis and internship experiences taught me how to analyze scientific work that is related to my research, learn from other researchers' experiences and place my work among them. This is crucial for analytical and exploratory skills that can be applied in both scientific and industrial research domains.

Bibliography

- [1] A. Azaria, “ChatGPT Usage and Limitations.” working paper or preprint, Dec. 2022.
- [2] R. Dale, “Gpt-3: What’s it good for?,” *Natural Language Engineering*, vol. 27, no. 1, p. 113–118, 2021.
- [3] L. Floridi and M. Chiriatti, “Gpt-3: Its nature, scope, limits, and consequences,” *Minds and Machines*, vol. 30, p. 681–694, Nov. 2020.
- [4] A. Hörnemalm, *ChatGPT as a Software Development Tool: The Future of Development*. PhD thesis, 2023.
- [5] J. K. Das, S. Mondal, and C. K. Roy, “Investigating the utility of chatgpt in the issue tracking system: An exploratory study,” 2024.
- [6] J. Liu, C. S. Xia, Y. Wang, and L. ZHANG, “Is your code generated by chatgpt really correct? rigorous evaluation of large language models for code generation,” in *Advances in Neural Information Processing Systems* (A. Oh, T. Neumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, eds.), vol. 36, pp. 21558–21572, Curran Associates, Inc., 2023.
- [7] Q. Zhang, T. Zhang, J. Zhai, C. Fang, B. Yu, W. Sun, and Z. Chen, “A critical review of large language model on software engineering: An example from chatgpt and automated program repair,” 2023.
- [8] A. Ahmad, M. Waseem, P. Liang, M. Fahmideh, M. S. Aktar, and T. Mikkonen, “Towards human-bot collaborative software architecting with chatgpt,” in *Proceedings of the 27th International Conference on Evaluation and Assessment in Software Engineering*, EASE ’23, (New York, NY, USA), p. 279–285, Association for Computing Machinery, 2023.
- [9] T. Xiao, C. Treude, H. Hata, and K. Matsumoto, “Devsgpt: Studying developer-chatgpt conversations,” Nov. 2023.
- [10] K. I. Roumeliotis and N. D. Tselikas, “Chatgpt and open-ai models: A preliminary review,” *Future Internet*, vol. 15, no. 6, 2023.

- [11] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. u. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Advances in Neural Information Processing Systems* (I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, eds.), vol. 30, Curran Associates, Inc., 2017.
- [12] M. L. Siddiq and J. C. Santos, “Generate and pray: Using salms to evaluate the security of llm generated code,” *arXiv preprint arXiv:2311.00889*, 2023.
- [13] M. F. Rabbi, A. I. Champa, M. F. Zibran, and M. R. Islam, “Ai writes, we analyze: The chatgpt python code saga,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 177–181, Association for Computing Machinery, 2024.
- [14] Y. Zhang, R. Meredith, W. Reeves, J. Coriolano, M. A. Babar, and A. Rahman, “Does generative ai generate smells related to container orchestration?: An exploratory study with kubernetes manifests,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 192–196, Association for Computing Machinery, 2024.
- [15] A. Clark, D. Igbokwe, S. Ross, and M. F. Zibran, “A quantitative analysis of quality and consistency in ai-generated code,” in *2024 7th International Conference on Software and System Engineering (ICoSSE)*, pp. 37–41, 2024.
- [16] K. Moratis, T. Diamantopoulos, D.-N. Nastos, and A. Symeonidis, “Write me this code: An analysis of chatgpt quality for producing source code,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 147–151, Association for Computing Machinery, 2024.
- [17] M. L. Siddiq, L. Roney, J. Zhang, and J. C. D. S. Santos, “Quality assessment of chatgpt generated code and their use by developers,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 152–156, Association for Computing Machinery, 2024.
- [18] B. Grewal, W. Lu, S. Nadi, and C.-P. Bezemer, “Analyzing developer use of chatgpt generated code in open source github projects,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 157–161, Association for Computing Machinery, 2024.

- [19] K. Jin, C.-Y. Wang, H. V. Pham, and H. Hemmati, “Can chatgpt support developers? an empirical evaluation of large language models for code generation,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 167–171, Association for Computing Machinery, 2024.
- [20] O. S. Chavan, D. D. Hinge, S. S. Deo, Y. O. Wang, and M. W. Mkaouer, “Analyzing developer-chatgpt conversations for software refactoring: An exploratory study,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 207–211, Association for Computing Machinery, 2024.
- [21] E. A. AlOmar, A. Venkatakrishnan, M. W. Mkaouer, C. Newman, and A. Ouni, “How to refactor this code? an exploratory study on developer-chatgpt refactoring conversations,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 202–206, Association for Computing Machinery, 2024.
- [22] S. Mondal, S. D. Bappon, and C. K. Roy, “Enhancing user interaction in chatgpt: Characterizing and consolidating multiple prompts for issue resolution,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 222–226, Association for Computing Machinery, 2024.
- [23] L. Wu, Y. Zhao, X. Hou, T. Liu, and H. Wang, “Chatgpt chats decoded: Uncovering prompt patterns for superior solutions in software development lifecycle,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 142–146, Association for Computing Machinery, 2024.
- [24] L. Aguiar, M. Paixao, R. Carmo, E. Soares, A. Leal, M. Freitas, and E. Gama, “Multi-language software development in the llm era: Insights from practitioners’ conversations with chatgpt,” in *Proceedings of the 18th ACM/IEEE International Symposium on Empirical Software Engineering and Measurement*, ESEM ’24, (New York, NY, USA), p. 489–495, Association for Computing Machinery, 2024.
- [25] M. Chouchen, N. Bessghaier, M. Begoug, A. Ouni, E. Alomar, and M. W. Mkaouer, “How do software developers use chatgpt? an exploratory study on github pull requests,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 212–216, Association for Computing Machinery, 2024.
- [26] M. Watanabe, Y. Kashiwa, B. Lin, T. Hirao, K. Yamaguchi, and H. Iida, “On the use of chatgpt for code review: Do developers like

- reviews by chatgpt?,” EASE ’24, (New York, NY, USA), p. 375–380, Association for Computing Machinery, 2024.
- [27] H. Hao and Y. Tian, “Engaging with ai: An exploratory study on developers’ sharing and reactions to chatgpt in github pull requests,” in *Proceedings of the 39th IEEE/ACM International Conference on Automated Software Engineering Workshops*, ASEW ’24, (New York, NY, USA), p. 156–160, Association for Computing Machinery, 2024.
 - [28] H. Hao, K. A. Hasan, H. Qin, M. Macedo, Y. Tian, S. H. H. Ding, and A. E. Hassan, “An empirical study on developers’ shared conversations with chatgpt in github pull requests and issues,” *Empirical Software Engineering*, vol. 29, p. 150, Sep 2024.
 - [29] E. Sagdic, A. Bayram, and M. R. Islam, “On the taxonomy of developers’ discussion topics with chatgpt,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 197–201, Association for Computing Machinery, 2024.
 - [30] M. Grootendorst, “Bertopic: Neural topic modeling with a class-based tf-idf procedure,” 2022.
 - [31] S. Mohamed, A. Parvin, and E. Parra, “Chatting with ai: Deciphering developer conversations with chatgpt,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 187–191, Association for Computing Machinery, 2024.
 - [32] A. I. Champa, M. F. Rabbi, C. Nachuma, and M. F. Zibran, “Chatgpt in action: Analyzing its use in software development,” in *Proceedings of the 21st International Conference on Mining Software Repositories*, MSR ’24, (New York, NY, USA), p. 182–186, Association for Computing Machinery, 2024.
 - [33] M. Lewis, Y. Liu, N. Goyal, M. Ghazvininejad, A. Mohamed, O. Levy, V. Stoyanov, and L. Zettlemoyer, “Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension,” 2019.
 - [34] L. Tan, “Pywsd: Python implementations of word sense disambiguation (wsd) technologies [software].” <https://github.com/alvations/pywsd>, 2014.
 - [35] S. Bird, E. Klein, and E. Loper, *Natural language processing with Python: analyzing text with the natural language toolkit*. ” O’Reilly Media, Inc.”, 2009.

- [36] Z. Elberrichi and B. Aljohar, “N-grams in texts categorization,” *Scientific Journal of King Faisal University (Basic and Applied Sciences)*, vol. 8, no. 2, pp. 25–39, 2007.
- [37] A. C. Mueller, “Wordcloud,” 2023.
- [38] C. Bou Rjeily, G. Badr, A. Hajjarm El Hassani, and E. Andres, “Medical data mining for heart diseases and the future of sequential mining in medical field,” in *Machine learning paradigms: Advances in data analytics*, pp. 71–99, Springer, 2018.
- [39] R. Agrawal and R. Srikant, “Mining sequential patterns,” in *Proceedings of the eleventh international conference on data engineering*, pp. 3–14, IEEE, 1995.
- [40] B. Requena, G. Cassani, J. Tagliabue, C. Greco, and L. Lacasa, “Shopper intent prediction from clickstream e-commerce data with minimal browsing information,” *Scientific reports*, vol. 10, no. 1, p. 16983, 2020.
- [41] S. Kadioğlu, X. Wang, A. Hosseininasab, and W.-J. van Hoeve, “Seq2pat: Sequence-to-pattern generation to bridge pattern mining with machine learning,” *AI Magazine*, vol. 44, no. 1, pp. 54–66, 2023.
- [42] X. Wang, A. Hosseininasab, P. Colunga, S. Kadioğlu, and W.-J. van Hoeve, “Seq2pat: Sequence-to-pattern generation for constraint-based sequential pattern mining,” *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 36, pp. 12665–12671, Jun. 2022.
- [43] B. Grün and K. Hornik, “topicmodels: An r package for fitting topic models,” *Journal of Statistical Software*, vol. 40, no. 13, p. 1–30, 2011.
- [44] Z. Tong and H. Zhang, “A text mining research based on lda topic modelling,” in *International conference on computer science, engineering and information technology*, pp. 201–210, 2016.
- [45] S. Syed and M. Spruit, “Full-text or abstract? examining topic coherence scores using latent dirichlet allocation,” in *2017 IEEE International Conference on Data Science and Advanced Analytics (DSAA)*, pp. 165–174, 2017.
- [46] A. Neishabouri, M. C. Desmarais, and P. Montreal, “Reliability of perplexity to find number of latent topics,” in *FLAIRS*, pp. 246–251, 2020.
- [47] “CountVectorizer — scikit-learn.org.” https://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html. [Accessed 23-09-2024].

- [48] “LatentDirichletAllocation — scikit-learn.org.” <https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.LatentDirichletAllocation.html>. [Accessed 23-09-2024].

Glossary

ChatGPT Chat Generative Pre-trained Transformer. 1–9, 12, 14–17, 22, 26–32

LDA Latent Dirichlet Allocation. 20, 26

LLM Large Language Model. 2, 30

NLP Natural Language Processing. 4, 7

SPM Sequential pattern mining. 1, 3, 17, 19, 24, 29, 30


Appendix A



Appendix

A.1 N-grams WordClouds



Figure A.1: N-grams from conversations sourced from GitHub commits.








Figure A.6: N-grams from conversations sourced from GitHub repositories.




Figure A.7: N-grams from conversations sourced from Hacker News.

A.2 Topic models

Topic name	Top words
Task-based file management	file, task, need, set, solve, solve task, use, working set, working, create
Function testing	test, use, file, name, code, find, std, return, gettile, tile
UI layout planning	plan, file, height, use, create, div, script, goal, go, import
AI-assisted development	file, task, use, junior, prompt, ai, project, create, work, script

Table A.1: Topic models with top words for conversations sourced from GitHub commits.

Topic name	Top words
Code & model behavior	use, file, th, return, line, code, model, input, error, string
Data & value processing	use, file, return, show, name, process, value, number, function, string

Table A.2: Topic models with top words for conversations sourced from GitHub issues.

Topic name	Top words
Product design feasibility	desirability, feasibility, customer, episode, technology, operational, technology feasibility, customer desirability, employee desirability, employee
Silicon Valley series	richard, pied, pied piper, piper, hi, use, hooli, file, erlich
Data handling	object, use, string, try short string, environment, email, want, agent, json, data
Node installation issues	new, node, value, error, please, install, npm, run, path
Change history	change, file, type, use, team, write, history, device, change history

Silicon Valley cast	hi, richard, belson, erlich, piper, pied piper, app, dinesh, gilfoyle
Messages & Actions	richard, say, return, hi, message, use, action, code, student, pied
Silicon Valley series: Pied Piper	richard, pied, pied piper, piper, use, hi, hooli, clause, service
Model calibration	agent, tool, model, policy, use, experiment, calibration

Table A.3: Topic models with top words for conversations sourced from GitHub discussions.

Topic name	Top words
C++ function signatures	return, std, string, use, value, int, vector, const, line, file
Node package metadata	use, author, node, npm, list, resource, id, code, file, return
Ontology & classification	subclassof, id, string, cavity, entity, return, book, genre, part, use
Mathematical constants & Types	const, sum, code, int, use, false, list, float, function, add
Web attestation APIs	string, use, user, code, return, attestation, like, web, request, attester
Item handling logic	item, use, string, want, code, need, file, function, please, return
Python packaging errors	program, pip, use, write, package, error, user, go, memory, model
Enterprise service logic	service, company, code, use, make, information, std, go, provide, get
File-based testing	use, test, file, add, run, function, image, style, echo, code
Module initialization errors	new, const, code, error, import, string, id, return, public, list

Table A.4: Topic models with top words for conversations sourced from GitHub pull requests.

Topic name	Top words
Step-by-step coding guidance	say, code, step, please, std, style, solution, end, use, detail
Evidence-supported argument	agent, argument, please, topic, subject, student, evidence, make, workshop, debate
Book passage creation	like, passage, use, make, hi, would, one, know, provide, time
OWL/RDF ontologies	owl, type, rdfs, rdf, rdf type, class, rdf type owl, type owl, relationship, ontology
Functional file operations	file, use, return, function, code, data, new, create, int, error

Table A.5: Topic models with top words for conversations sourced from GitHub code files

Topic name	Top words
UI layout metadata	height, metadata, list, skyscraper, viewport, fill, height fill, data, employee, salary
Research effort	planet, help, disaster, increase, research, program, also, reduce, effort, data

Table A.6: Topic models with top words for conversations sourced from GitHub repositories.

Topic name	Top words
Conversational requests	could, would, hi, make, go, time, question, one, use, like
User-host session actions	make, user, host, write, user host, word, goal, please, get, like
Narrative creation	one, story, use, like, hi, would, give, create, question, think
Game interaction errors	one, give, like, game, use, add, code, medical, error, would
File tracking utilities	use, track, let, file, return, like, one, time, point, value

Company schedule	company, company company, company company company, company company company company, headline, schedule, insert, level, like, tell
Positional pattern capture	position, position position, position position po- sition, like, capture, capture position, capture position position, capture position position po- sition, line, form
Code support	use, make, get, time, work, one, code, like, need, give

Table A.7: Topic models with top words for conversations sourced from HackerNews.

A.3 Topic models' WordClouds




Figure A.8: Topic models for conversations sourced from GitHub commits.




Figure A.9: Topic models for conversations sourced from GitHub issues.




Figure A.11: Topic models for conversations sourced from GitHub pull requests.



Figure A.12: Topic models for conversations sourced from GitHub code files.



Figure A.13: Topic models for conversations sourced from GitHub repositories.

